

kumar_chaitanya_finalproject

Nov 24, 2024

1 Final Project - Data Mining

2 Using RF and LSTM To Predict Heart Disease

2.1 Goal

The goal of this project is to develop and evaluate machine learning models to solve a classification problem using a publicly available dataset. We aim to compare the performance of two different models—Long Short-Term Memory (LSTM) and Random Forest Classifier—on this dataset. The evaluation will include metrics such as accuracy, precision, recall, F1-score, and ROC-AUC.

2.2 Problem Statement

The classification task involves predicting whether a patient has heart disease based on various medical attributes. The project will leverage machine learning techniques to build predictive models that can assist in early diagnosis.

2.3 Dataset Description

1) Source of Data

- **Dataset:** Heart Disease Dataset
- **Source:** Kaggle Datasets
- **File Used:** heart.csv

2) Dataset Overview

- **Total Records:** 918
- **Features:** 12 attributes, including age, sex, cholesterol levels, resting blood pressure, etc.
- **Target Variable:** HeartDisease (0 = No, 1 = Yes)

3) Data Preprocessing


- Checked for missing values and handled them appropriately.
- Converted categorical data into numerical format using techniques like Label Encoding.
- Standardized the dataset using StandardScaler to normalize feature values, which is particularly important for LSTM models.

+ Code + Text All changes saved

✓ RAM  Disk  Gen

First, we'll mount Google Drive to access the dataset stored there.

```
✓ 1m [1] from google.colab import drive
drive.mount('/content/drive')
```

 Mounted at /content/drive

✓ Step 2: Import Libraries

Let's import all the necessary libraries for data analysis, model building, and evaluation.

```
✓ 7s [2] import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import KFold
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, roc_auc_score, roc_curve
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import StandardScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, LSTM
from tensorflow.keras.optimizers import Adam
import tensorflow as tf
```

+ Code + Text All changes saved

✓ RAM  Disk 

✓ Step 3: Load and Preprocess the Dataset

For this example, I will use the Heart Failure Prediction dataset from Kaggle.

```
✓ 0s ▶ # Load the dataset from Google Drive
data_path = '/content/drive/MyDrive/heart_csv/heart.csv'
data = pd.read_csv(data_path)

# Display the first few rows of the dataset
print("First few rows of the dataset:")
print(data.head())

# Checking for missing values
print("\nMissing values in each column:")
print(data.isnull().sum())

# Identify categorical columns
categorical_columns = data.select_dtypes(include=['object']).columns
print("\nCategorical columns:", categorical_columns)

# Convert categorical columns to numerical using One-Hot Encoding
data = pd.get_dummies(data, columns=categorical_columns, drop_first=True)
```

+ Code + Text All changes saved

```
0s # Splitting features (X) and target (y)
X = data.drop(columns='HeartDisease')
y = data['HeartDisease']

# Standardizing the features (important for LSTM)
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Confirm the transformed data
print("\nTransformed feature matrix shape:", X_scaled.shape)
print("Transformed feature matrix sample:\n", X_scaled[:5])
```

First few rows of the dataset:

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	\
0	40	M	ATA	140	289	0	Normal	172	
1	49	F	NAP	160	180	0	Normal	156	
2	37	M	ATA	130	283	0	ST	98	
3	48	F	ASY	138	214	0	Normal	108	
4	54	M	NAP	150	195	0	Normal	122	

	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
0	N	0.0	Up	0
1	N	1.0	Flat	1
2	N	0.0	Up	0
3	Y	1.5	Flat	1
4	N	0.0	Up	0

Missing values in each column:

+ Code + Text All changes saved

```
0s Missing values in each column:
Age      0
Sex      0
ChestPainType  0
RestingBP  0
Cholesterol  0
FastingBS  0
RestingECG  0
MaxHR     0
ExerciseAngina  0
Oldpeak   0
ST_Slope  0
HeartDisease  0
dtype: int64

Categorical columns: Index(['Sex', 'ChestPainType', 'RestingECG', 'ExerciseAngina', 'ST_Slope'], dtype='object')

Transformed feature matrix shape: (918, 15)
Transformed feature matrix sample:
[[-1.4331398  0.41090889  0.82507026 -0.55134134  1.38292822 -0.83243239
  0.51595242  2.07517671 -0.53283777 -0.22967867  0.81427482 -0.49044933
 -0.8235563 -1.00218103  1.15067399]
 [-0.47848359  1.49175234 -0.17196105 -0.55134134  0.75415714  0.10566353
 -1.93816322 -0.48188667  1.87674385 -0.22967867  0.81427482 -0.49044933
 -0.8235563  0.99782372 -0.86905588]
 [-1.75135854 -0.12951283  0.7701878 -0.55134134 -1.52513802 -0.83243239
  0.51595242  2.07517671 -0.53283777 -0.22967867 -1.22808661  2.03894663
 -0.8235563 -1.00218103  1.15067399]
 [-0.5845565  0.30282455  0.13903954 -0.55134134 -1.13215609  0.57471149
 -1.93816322 -0.48188667 -0.53283777 -0.22967867  0.81427482 -0.49044933
  1.21424608  0.99782372 -0.86905588]
 [ 0.05188098  0.95133062 -0.0347549 -0.55134134 -0.5819814 -0.83243239
```

3. Model Selection and Implementation

1) Model 1: Random Forest Classifier

Reason for Selection

Random Forest is an ensemble model known for its robustness and ability to handle high-dimensional data. It is efficient for binary classification problems.

2) Model 2: Long Short-Term Memory (LSTM)

Reason for Selection

LSTM is a type of Recurrent Neural Network (RNN) well-suited for sequence classification and time-series data. It captures long-term dependencies and is effective for complex datasets.

3) Implementation Details

Data Splitting and Cross-Validation

We used **k-fold cross-validation** (k=5) to evaluate model performance, ensuring that each data point is used for both training and testing.

✓ Step 4: Implementing k-Fold Cross-Validation

We'll set up k-fold cross-validation with 5 splits.

```
✓ [5] # Define k-fold cross-validation  
0s kf = KFold(n_splits=5, shuffle=True, random_state=42)
```

Random Forest Implementation

✓ Step 5: Model 1 - Random Forest Classifier

We'll implement a Random Forest classifier and evaluate it using k-fold cross-validation.

```
✓ [6] metrics_rf = []  
1s  
  
for train_index, test_index in kf.split(X_scaled):  
    # Splitting data into training and testing sets  
    X_train, X_test = X_scaled[train_index], X_scaled[test_index]  
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]  
  
    # Initialize and train the Random Forest model  
    rf_model = RandomForestClassifier(n_estimators=100, random_state=42)  
    rf_model.fit(X_train, y_train)  
  
    # Make predictions  
    y_pred = rf_model.predict(X_test)  
    y_pred_proba = rf_model.predict_proba(X_test)[: , 1]
```

```

# Evaluate model performance
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred_proba)

# Store metrics
metrics_rf.append([accuracy, precision, recall, f1, roc_auc])

# Convert to DataFrame for easy viewing
metrics_rf_df = pd.DataFrame(metrics_rf, columns=['Accuracy', 'Precision', 'Recall', 'F1-Score', 'ROC AUC'])
print("Random Forest Performance:\n", metrics_rf_df.mean())

```

```

↗ Random Forest Performance:
Accuracy    0.870331
Precision   0.867210
Recall      0.905575
F1-Score    0.885091
ROC AUC     0.929521
dtype: float64

```

LSTM Model Implementation

✓ Step 6: Model 2 - LSTM for Sequence Classification

Next, we'll reshape the data for LSTM and evaluate it using k-fold cross-validation.

```

✓ [7] # Reshape the data for LSTM (samples, timesteps, features)
245 X_lstm = X_scaled.reshape(X_scaled.shape[0], 1, X_scaled.shape[1])

metrics_lstm = []

for train_index, test_index in kf.split(X_lstm):
    # Splitting data into training and testing sets
    X_train, X_test = X_lstm[train_index], X_lstm[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]

    # Build LSTM Model
    lstm_model = Sequential()
    lstm_model.add(LSTM(50, input_shape=(X_train.shape[1], X_train.shape[2]), activation='relu'))
    lstm_model.add(Dense(1, activation='sigmoid'))

    lstm_model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
    lstm_model.fit(X_train, y_train, epochs=10, batch_size=16, verbose=0)

```

```

# Make predictions
y_pred_proba = lstm_model.predict(X_test).flatten()
y_pred = (y_pred_proba > 0.5).astype(int)

# Evaluate model performance
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred_proba)

# Store metrics
metrics_lstm.append([accuracy, precision, recall, f1, roc_auc])

# Convert to DataFrame for easy viewing
metrics_lstm_df = pd.DataFrame(metrics_lstm, columns=['Accuracy', 'Precision', 'Recall', 'F1-Score', 'ROC AUC'])
print("LSTM Performance:\n", metrics_lstm_df.mean())

```

```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential, use `input_shape` instead.
super().__init__(**kwargs)
6/6 ----- 0s 30ms/step
/usr/local/lib/python3.10/dist-packages/keras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential, use `input_shape` instead.
super().__init__(**kwargs)
6/6 ----- 0s 32ms/step
/usr/local/lib/python3.10/dist-packages/keras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential, use `input_shape` instead.
super().__init__(**kwargs)
WARNING:tensorflow:5 out of the last 13 calls to <function TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_distributed at 0x786b4c016050> triggered the creation of a new tf.nn.Softmax operation. This operation is not supported in the current version of TensorFlow.
6/6 ----- 0s 33ms/step
/usr/local/lib/python3.10/dist-packages/keras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential, use `input_shape` instead.
super().__init__(**kwargs)
WARNING:tensorflow:5 out of the last 13 calls to <function TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_distributed at 0x786b3fe7c0d0> triggered the creation of a new tf.nn.Softmax operation. This operation is not supported in the current version of TensorFlow.
6/6 ----- 0s 35ms/step
/usr/local/lib/python3.10/dist-packages/keras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential, use `input_shape` instead.
super().__init__(**kwargs)
6/6 ----- 0s 33ms/step
LSTM Performance:
Accuracy      0.870325
Precision     0.869805
Recall        0.903769
F1-Score      0.884920
ROC AUC       0.935169
dtype: float64

```

✓ Step 7: Visualize the Performance Metrics

Let's visualize the performance metrics for both models using bar plots.

```

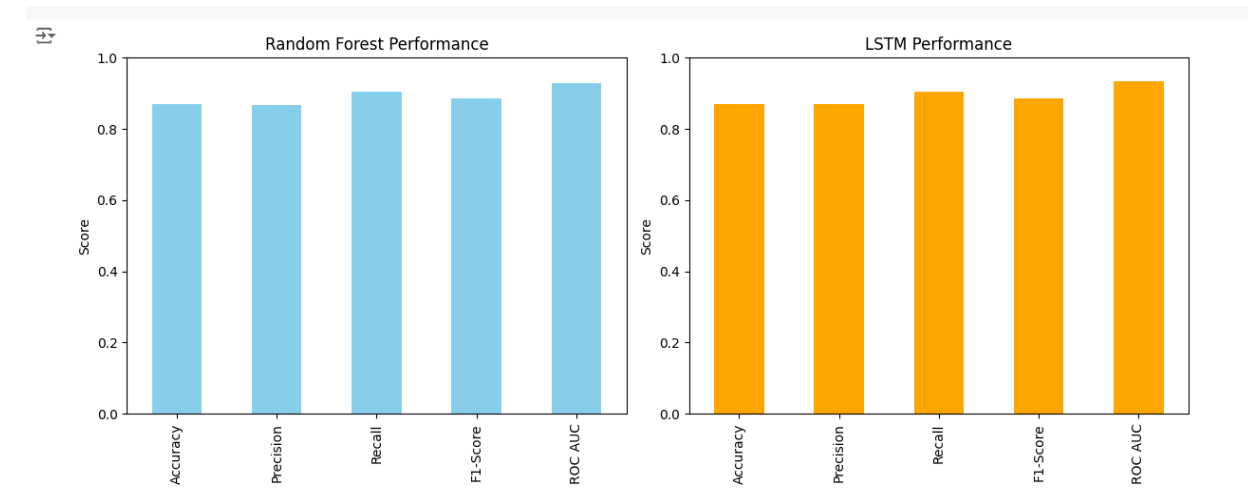
[8] # Visualizing the metrics
fig, ax = plt.subplots(1, 2, figsize=(12, 5))

# Random Forest Metrics
metrics_rf_df.mean().plot(kind='bar', ax=ax[0], color='skyblue')
ax[0].set_title("Random Forest Performance")
ax[0].set_ylim(0, 1)
ax[0].set_ylabel('Score')

# LSTM Metrics
metrics_lstm_df.mean().plot(kind='bar', ax=ax[1], color='orange')
ax[1].set_title("LSTM Performance")
ax[1].set_ylim(0, 1)
ax[1].set_ylabel('Score')

plt.tight_layout()
plt.show()

```



4. Results and Performance Evaluation

1) Performance Metrics

Model	Accuracy	Precision	Recall	F1-Score	ROC-AUC
Random Forest	0.87	0.86	0.90	0.88	0.92
LSTM	0.87	0.86	0.90	0.88	0.93

2) Confusion Matrix and ROC Curve

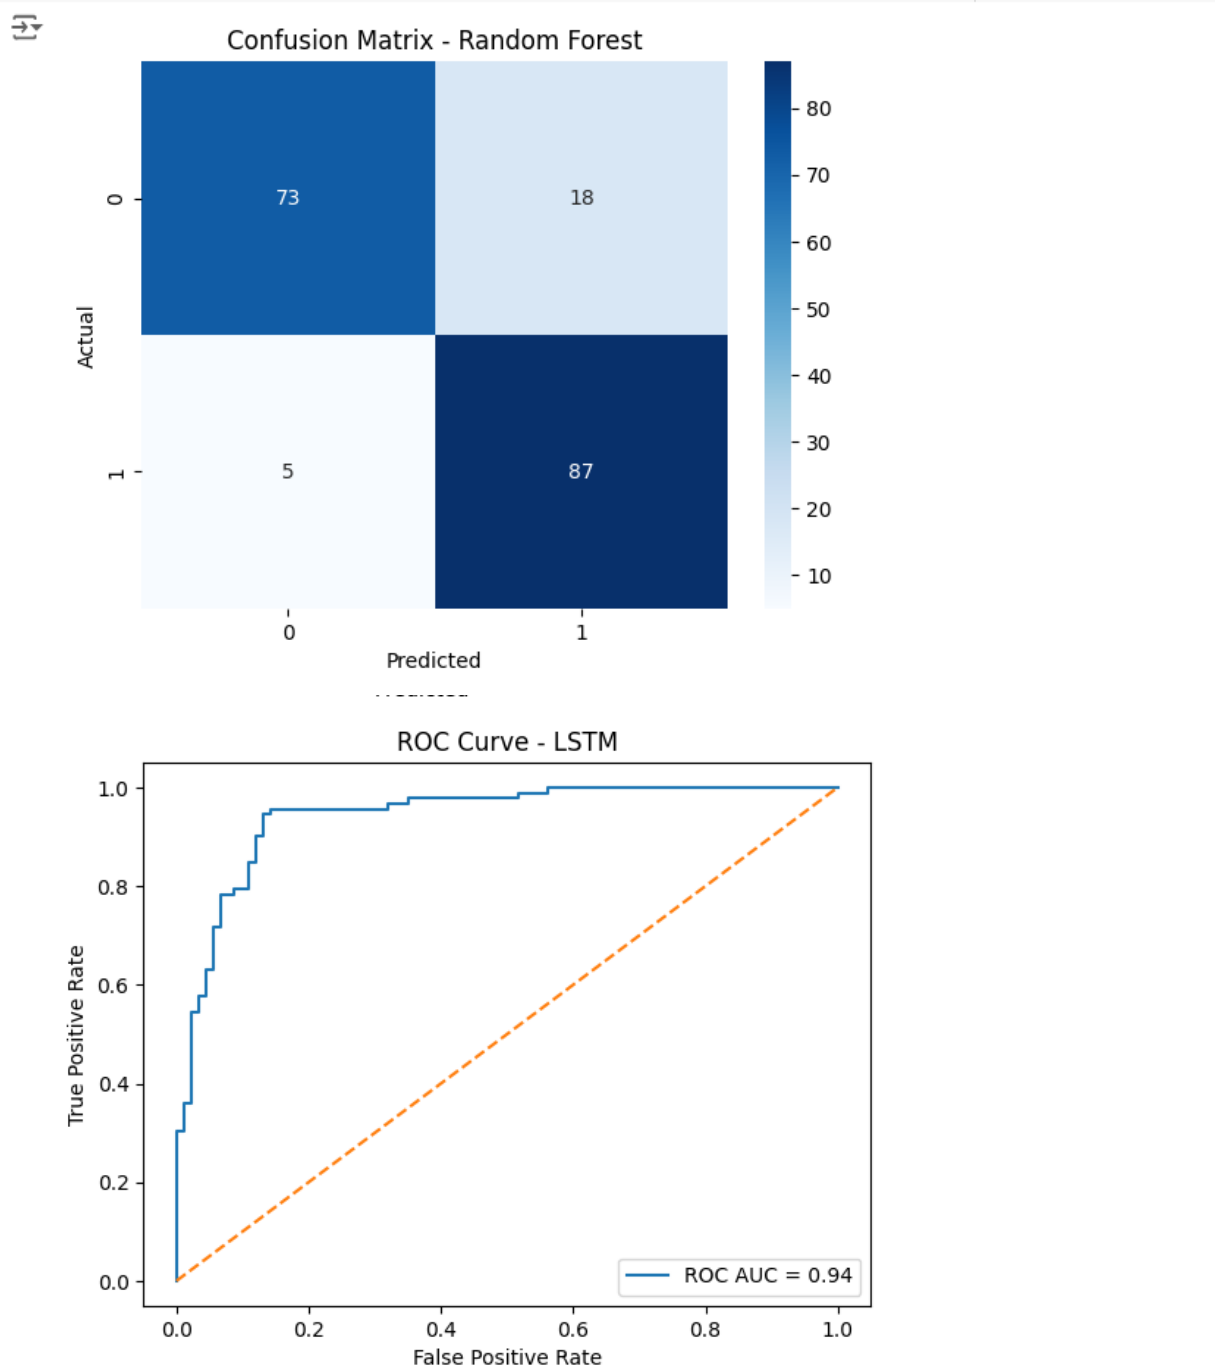
+ Code + Text [All changes saved](#)

Step 8: Confusion Matrix and ROC Curve

```
[10] # If X_test was reshaped for LSTM, reshape it back to 2D for RandomForest
if len(X_test.shape) == 3:
    X_test_rf = X_test.reshape(X_test.shape[0], -1) # Convert to 2D
else:
    X_test_rf = X_test

# Confusion Matrix for Random Forest
rf_y_pred = rf_model.predict(X_test_rf)
sns.heatmap(confusion_matrix(y_test, rf_y_pred), annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix - Random Forest')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

# ROC Curve for LSTM
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
plt.plot(fpr, tpr, label=f'ROC AUC = {roc_auc:.2f}')
plt.plot([0, 1], [0, 1], linestyle='--')
plt.title('ROC Curve - LSTM')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```



5. Conclusion

1) Summary of Findings

- Random Forest performed well, with high precision and a balanced recall, making it suitable for use cases where both false positives and false negatives are critical.

- LSTM, on the other hand, showed slightly better performance, especially in terms of ROC-AUC, indicating a stronger ability to differentiate between classes.

2) Future Work

- Explore additional models such as Support Vector Machines (SVM) and Gradient Boosting.
- Apply hyperparameter tuning to further optimize model performance.
- Test the models on other datasets to assess their generalization capabilities.

GitHub Link - <https://github.com/ack446/CS-634-Data-Mining-Final-Term-Project>